


Agentic Reorganization Orchestration

Turn a freeform reorg request into a governed, human-approved, auditable set of changes across HR, Finance, and Legal.

AI proposes	Policy decides	Humans approve	Tools execute
-------------	----------------	----------------	---------------



1. Assumptions and approval policy

Important: These are illustrative design assumptions, not universal HR, Finance, or Legal rules. Policy owners must validate them by change type, threshold, business unit, and jurisdiction.

Approval rules

HR approves every employee or org-structure change.
Finance approves headcount, budget, cost-center, or GL impact.
Legal is required for RIF/position elimination and cross-country moves; conditional for employment-term or consultation risk.
No self-approval for high-risk cases.

Data and systems

HRIS owns employee, position, manager, and org data.
Planning/ERP owns budget, cost-center, and GL data.
Stable IDs exist; some steps have no API.
Downstream state supports read-back.

Execution and AI


Changes are effective-dated and support dry run.
HRIS precedes Planning, Finance, and IAM unless policy differs.
AI recommends; it does not set policy or make legal judgments.
Ambiguity triggers human clarification.

Illustrative approval matrix

Change	HR	Finance	Legal
Manager only	Required	No	Usually no
Team transfer	Required	Conditional	Conditional
Cost center / budget	Required	Required	Usually no
Role / location / terms	Required	Conditional	Risk-based
Cross-country / RIF	Required	Required	Required

2. End-to-end technical architecture

The model interprets intent; deterministic services control authorization and side effects. One correlation ID links the source message to every decision, tool call, and read-back result.



1. Ingest	Authenticate requester, scan attachments, preserve source, assign correlation ID.
2. Interpret	Super Agent emits a typed routing envelope; low confidence stops for clarification.
3. Build case	Resolve names to stable IDs and create typed, sourced before/after changes.
4. Govern	Rules validate data, derive risk, select gates, and bind approval to the case hash.
5. Execute	Workflow runs the approved dependency graph through narrow connectors or human tasks.
6. Reconcile	Read-back compares observed state with approved intent; mismatch opens recovery.

3. Worked example

Move the Analytics team from Marketing to Finance

Incoming request

“Move the Analytics team from Marketing to Finance under Jane Smith on Oct 1. Transfer the team budget too. Two employees are in Germany.”

Super Agent output

Domain: HR_REORG
Intent: TEAM_TRANSFER
Risk: budget + multi-jurisdiction
Confidence: 0.94
Route: Reorg/HR Agent

Canonical case

REORG-2026-001 / v1
Employees: 12
Marketing -> Finance
Manager: Jane Smith
Effective: Oct 1
Countries: US, Germany

Approval route

HR: required
Finance: required
Legal: required by illustrative
Germany consultation-risk
trigger

Approvals bind to case v1 hash.

Execution journey



Assumptions applied

HR: employees, manager, and org structure change. **Finance:** budget and cost center move. **Legal:** Germany is an illustrative consultation-risk signal; policy triggers review and Legal decides. **Order:** HRIS -> Planning/Finance -> IAM is illustrative.

4. Operating model

Canonical schema

Identity: reorg_id, case_version, status
Source: channel, requester, timestamp, reference
Scope: employees, teams, countries, effective date
Change set: typed before/after values + provenance
Risk: policy signals and confidence
Approvals: role, actor, case hash, policy version
Plan: dependency steps and idempotency keys
Reconciliation: expected vs observed state

Lifecycle

RECEIVED -> CLASSIFIED -> VALIDATED ->
PLANNED -> APPROVED -> SCHEDULED ->
EXECUTING -> RECONCILING -> COMPLETED

Exception states:

NEEDS CLARIFICATION / PAUSED / PARTIAL /
FAILED

Material edits create a new case version and
invalidate prior approvals.

Dependency model

Prechecks -> HRIS -> Planning -> Finance -> IAM -> Read-back

Each step declares prerequisites, owner, timeout, retry policy, idempotency key, expected result, compensation, and required evidence. Independent steps may run in parallel; failed prerequisites block dependents.

Human-in-the-loop and no-API work

- Humans clarify ambiguous scope and approve the exact impact preview.
- Manual tasks have a named owner, due date, minimum-PII instructions, and evidence requirement.
- Attestation alone is insufficient: verify completion through read-back when possible.
- Operators can pause, retry, compensate, or create a remediation case.

5. Trust, resilience, and audit

Security and PII

SSO + MFA; RBAC/ABAC; separation of duties; encryption; DLP; retention and residency controls; redact/tokenize before model calls; treat attachments as untrusted.

Failure and recovery

Idempotent retry; resume from checkpoint; safe compensation; roll-forward remediation; freeze related cases and escalate when consistency is uncertain.

Observability

Track latency, clarification, correction, denial, tool error, retry, duplicate-prevention, and mismatch rates. Alert on SLA breaches and policy bypass attempts.

Blast radius

Dry run by default; canary employees; per-case and batch limits; reversible steps first; connector circuit breakers; global kill switch.

Auditability

Record actor, case hash, policy/model/tool version, before/after snapshot, approval evidence, tool result, and reconciliation evidence under one correlation ID.

Consistency model

This is a saga, not a cross-system transaction. "Completed" means observed downstream state matches approved intent - not merely that API calls succeeded.

6. Prototype and production roadmap

Prototype - one scenario

Analytics-team transfer; synthetic data; HR + Finance approvals; Legal trigger; mocked connectors; one manual step; failure injection; audit export.

Pilot - one geography

Read-only shadow first; limited HRIS writes; canary batches; live reconciliation; operations console; incident playbooks.

Production - expand by evidence

Add jurisdictions and change types only after evaluation, policy ownership, connector certification, recovery drills, and privacy/security sign-off.

Prototype success measures

At least 95% correct routing on curated tests | **100%** of writes tied to valid approvals | **0** duplicate effects on replay

Alternatives considered

Option	Why not primary	What we retain
One general agent	Too much privilege; hard to assure	Natural-language intake
Forms + deterministic workflow	Poor with ambiguity and exceptions	Policy and execution control
Point-to-point RPA	Brittle and fragmented	Legacy adapters behind narrow tools
Distributed transaction	Impossible across SaaS and manual work	Saga, checkpoints, compensation

Open questions

What is authoritative by field? Who owns policy? What triggers Legal review by country? Which changes are never automated? What are batch limits, payroll/ledger windows, evidence retention, recovery objectives, and kill-switch authority?